

test case

September 24, 2022

```
[1]: import numpy as np
import cv2

# Open the video
cap = cv2.VideoCapture('C://Users/Nikolay/Downloads/video/Camera_
↳3_20220526_003249(2).mp4')

# Initialize frame counter
cnt = 1

w_frame, h_frame = int(cap.get(cv2.CAP_PROP_FRAME_WIDTH)), int(cap.get(cv2.
↳CAP_PROP_FRAME_HEIGHT))
fps, frames = cap.get(cv2.CAP_PROP_FPS), cap.get(cv2.CAP_PROP_FRAME_COUNT)

# define cropping values
x,y,h,w = 115,210,235,235

# output
fourcc = cv2.VideoWriter_fourcc(*'XVID')
out = cv2.VideoWriter('result.avi', fourcc, fps, (w, h))

while(cap.isOpened()):
    ret, frame = cap.read()

    cnt += 1 # Counting frames

    # Avoid problems when video finish
    if ret==True:
        # Cropping the frame
        crop_frame = frame[y:y+h, x:x+w]

        # Percentage
        xx = cnt *100/frames
        print(int(xx), '%')

    #Saving from the desired frames
```

```
cap.release()  
out.release()  
cv2.destroyAllWindows()
```

2

```
99 %
99 %
99 %
99 %
99 %
99 %
99 %
99 %
99 %
99 %
99 %
99 %
99 %
99 %
100 %
```

```
[7]: #size for reshape
```

```
down_width = 116
```

```
down_height = 116
```

```
down_points = (down_width, down_height)
```

```
[8]: # reshape all images to 116*116
```

```
vidcap = cv2.VideoCapture('C://Users/Nikolay/jupyter_notebooks/result.avi')
```

```
success,image = vidcap.read()
```

```
count = 0
```

```
while success:
```

```
    image = cv2.resize(image, down_points, interpolation= cv2.INTER_LINEAR)
```

```
    cv2.imwrite("frame%d.jpg" % count, image)      # save frame as JPEG file
```

```
↪
```

```
    success,image = vidcap.read()
```

```
    print('Read a new frame: ', success)
```

```
    count += 1
```

```
Read a new frame: True
```

```
Read a new frame: True
```

```
Read a new frame: True
```

```
Read a new frame: True
```

```
Read a new frame: True
```

```
Read a new frame: True
```

```
Read a new frame: True
```

```
Read a new frame: True
```

```
Read a new frame: True
```

```
Read a new frame: True
```

```
Read a new frame: True
```

```
Read a new frame: True
```

[illegible]

```
[16]: # checking the image
image = cv2.imread('C://Users/Nikolay/Desktop/test/1/frame9484.jpg')
cv2.imshow('Resized Down by defining height and width', image)
cv2.waitKey()
cv2.destroyAllWindows()
```

```
[3]: import torch
import random
#import numpy as np
import os

torch.backends.cudnn.deterministic = True
from torchvision.datasets import ImageFolder
from torchvision.transforms import ToTensor
from torchvision import datasets

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import Dataset
import torchvision.transforms as transforms
from torch.utils.data import DataLoader
import torchvision
from skimage import io
import pandas as pd

# fix random seed
random.seed(0)
#np.random.seed(0)
torch.manual_seed(0)
torch.cuda.manual_seed(0)

# using cuda
device = "cuda" if torch.cuda.is_available() else "cpu"
```

C:\Users\Nikolay\AppData\Local\Programs\Python\Python310\lib\site-packages\torch\random.py:42: UserWarning: Failed to initialize NumPy: module compiled against API version 0x10 but this version of numpy is 0xf (Triggered internally at ..\torch\src\utils\tensor_numpy.cpp:68.)
return default_generator.manual_seed(seed)

0.1 Making labels

```
[31]: X = os.listdir('C://Users/Nikolay/Desktop/test/no')
y = os.listdir('C://Users/Nikolay/Desktop/test/target')
```

```
[32]: print(X[0], y[0])
```

frame0.jpg frame10214.jpg

```
[33]: len(X), len(y)
```

```
[33]: (29105, 8578)
```

```
[54]: df_sit = pd.DataFrame({'Frame': y})  
df_Notsit = pd.DataFrame({'Frame': X})
```

```
[55]: df_sit['class'] = np.ones(len(y))  
df_Notsit['class'] = np.zeros(len(X))
```

```
[56]: df_Notsit.head()
```

```
[56]:
```

	Frame	class
0	frame0.jpg	0.0
1	frame1.jpg	0.0
2	frame10.jpg	0.0
3	frame100.jpg	0.0
4	frame1000.jpg	0.0

```
[59]: labels = pd.concat([df_sit, df_Notsit], axis = 0)  
labels.head()
```

```
[59]:
```

	Frame	class
0	frame10214.jpg	1.0
1	frame10215.jpg	1.0
2	frame10216.jpg	1.0
3	frame10217.jpg	1.0
4	frame10218.jpg	1.0

```
[136]: labels.to_excel('C://Users/Nikolay/jupyter_notebooks/labels.xlsx',header = True,  
                    index=['Frame', 'class'],  
                    columns=['Frame', 'class'])
```

```
[10]: pd.read_excel('C://Users/Nikolay/Desktop/test/images/labels.xlsx', index_col=0)
```

```
[10]:
```

	Frame	class
0	frame10214.jpg	1
1	frame10215.jpg	1
2	frame10216.jpg	1
3	frame10217.jpg	1
4	frame10218.jpg	1
...
29100	frame9995.jpg	0
29101	frame9996.jpg	0
29102	frame9997.jpg	0
29103	frame9998.jpg	0
29104	frame9999.jpg	0

```
[37683 rows x 2 columns]
```

0.2 making class for Dataset

```
[9]: class SitNotSitDataset(Dataset):
    def __init__(self, xlsx_file, root_dir, transform = None):
        self.annotations = pd.read_excel(xlsx_file, index_col=0)
        self.root_dir = root_dir
        self.transform = transform

    def __len__(self):
        return len(self.annotations)

    def __getitem__(self, index):
        img_path = os.path.join(self.root_dir, self.annotations.iloc[index, 0])
        image = io.imread(img_path)
        y_label = torch.tensor(int(self.annotations.iloc[index, 1]))

        if self.transform:
            image = self.transform(image)

        return(image, y_label)
```

```
[10]: dataset = SitNotSitDataset(xlsx_file = 'C://Users/Nikolay/Desktop/test/images/
↳labels.xlsx',
                                root_dir = 'C://Users/Nikolay/Desktop/test/images',
                                transform = transforms.ToTensor())
```

```
[11]: len(dataset)
```

```
[11]: 37683
```

0.3 train_test_split

```
[12]: train_size = int(0.7 * len(dataset))
test_size = len(dataset) - train_size
train_set, test_set = torch.utils.data.random_split(dataset, [train_size,
↳test_size])
```

```
[13]: batch_size = 16
```

```
[14]: train_loader = DataLoader(dataset = train_set, batch_size = batch_size, shuffle=
↳True)
test_loader = DataLoader(dataset = test_set, batch_size = batch_size, shuffle =
↳True)
```

0.4 using googlenet

```
[15]: model = torchvision.models.googlenet(pretrained = True)
```

```
C:\Users\Nikolay\AppData\Local\Programs\Python\Python310\lib\site-  
packages\torchvision\models\_utils.py:208: UserWarning: The parameter  
'pretrained' is deprecated since 0.13 and will be removed in 0.15, please use  
'weights' instead.  
  warnings.warn(  
C:\Users\Nikolay\AppData\Local\Programs\Python\Python310\lib\site-  
packages\torchvision\models\_utils.py:223: UserWarning: Arguments other than a  
weight enum or `None` for 'weights' are deprecated since 0.13 and will be  
removed in 0.15. The current behavior is equivalent to passing  
`weights=GoogLeNet_Weights.IMAGENET1K_V1`. You can also use  
`weights=GoogLeNet_Weights.DEFAULT` to get the most up-to-date weights.  
  warnings.warn(msg)
```

```
[18]: model.to(device)
```

```
[18]: GoogLeNet(  
  (conv1): BasicConv2d(  
    (conv): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3),  
    bias=False)  
    (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True,  
    track_running_stats=True)  
  )  
  (maxpool1): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,  
    ceil_mode=True)  
  (conv2): BasicConv2d(  
    (conv): Conv2d(64, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)  
    (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True,  
    track_running_stats=True)  
  )  
  (conv3): BasicConv2d(  
    (conv): Conv2d(64, 192, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),  
    bias=False)  
    (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,  
    track_running_stats=True)  
  )  
  (maxpool2): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,  
    ceil_mode=True)  
  (inception3a): Inception(  
    (branch1): BasicConv2d(  
      (conv): Conv2d(192, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)  
      (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True,  
    track_running_stats=True)  
    )  
    (branch2): Sequential(  

```



```

        (0): BasicConv2d(
          (conv): Conv2d(192, 96, kernel_size=(1, 1), stride=(1, 1), bias=False)
          (bn): BatchNorm2d(96, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
        (1): BasicConv2d(
          (conv): Conv2d(96, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
          (bn): BatchNorm2d(128, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
      (branch3): Sequential(
        (0): BasicConv2d(
          (conv): Conv2d(192, 16, kernel_size=(1, 1), stride=(1, 1), bias=False)
          (bn): BatchNorm2d(16, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
        (1): BasicConv2d(
          (conv): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
          (bn): BatchNorm2d(32, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
      (branch4): Sequential(
        (0): MaxPool2d(kernel_size=3, stride=1, padding=1, dilation=1,
ceil_mode=True)
        (1): BasicConv2d(
          (conv): Conv2d(192, 32, kernel_size=(1, 1), stride=(1, 1), bias=False)
          (bn): BatchNorm2d(32, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
    )
    (inception3b): Inception(
      (branch1): BasicConv2d(
        (conv): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(128, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
      )
      (branch2): Sequential(
        (0): BasicConv2d(
          (conv): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
          (bn): BatchNorm2d(128, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )

```

```

        (1): BasicConv2d(
          (conv): Conv2d(128, 192, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
          (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
    (branch3): Sequential(
      (0): BasicConv2d(
        (conv): Conv2d(256, 32, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(32, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
      )
      (1): BasicConv2d(
        (conv): Conv2d(32, 96, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn): BatchNorm2d(96, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    (branch4): Sequential(
      (0): MaxPool2d(kernel_size=3, stride=1, padding=1, dilation=1,
ceil_mode=True)
      (1): BasicConv2d(
        (conv): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
  )
  (maxpool3): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,
ceil_mode=True)
  (inception4a): Inception(
    (branch1): BasicConv2d(
      (conv): Conv2d(480, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch2): Sequential(
      (0): BasicConv2d(
        (conv): Conv2d(480, 96, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(96, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
      )
      (1): BasicConv2d(
        (conv): Conv2d(96, 208, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)

```

```

        (bn): BatchNorm2d(208, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (branch3): Sequential(
      (0): BasicConv2d(
        (conv): Conv2d(480, 16, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(16, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
      )
      (1): BasicConv2d(
        (conv): Conv2d(16, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn): BatchNorm2d(48, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    (branch4): Sequential(
      (0): MaxPool2d(kernel_size=3, stride=1, padding=1, dilation=1,
ceil_mode=True)
      (1): BasicConv2d(
        (conv): Conv2d(480, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    )
    (inception4b): Inception(
      (branch1): BasicConv2d(
        (conv): Conv2d(512, 160, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(160, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
      )
      (branch2): Sequential(
        (0): BasicConv2d(
          (conv): Conv2d(512, 112, kernel_size=(1, 1), stride=(1, 1), bias=False)
          (bn): BatchNorm2d(112, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
        (1): BasicConv2d(
          (conv): Conv2d(112, 224, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
          (bn): BatchNorm2d(224, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
    )
    (branch3): Sequential(

```

```

        (0): BasicConv2d(
          (conv): Conv2d(512, 24, kernel_size=(1, 1), stride=(1, 1), bias=False)
          (bn): BatchNorm2d(24, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
        (1): BasicConv2d(
          (conv): Conv2d(24, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
          (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
      (branch4): Sequential(
        (0): MaxPool2d(kernel_size=3, stride=1, padding=1, dilation=1,
ceil_mode=True)
        (1): BasicConv2d(
          (conv): Conv2d(512, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
          (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
    )
    (inception4c): Inception(
      (branch1): BasicConv2d(
        (conv): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(128, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
      )
      (branch2): Sequential(
        (0): BasicConv2d(
          (conv): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
          (bn): BatchNorm2d(128, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
        (1): BasicConv2d(
          (conv): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
          (bn): BatchNorm2d(256, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
      (branch3): Sequential(
        (0): BasicConv2d(
          (conv): Conv2d(512, 24, kernel_size=(1, 1), stride=(1, 1), bias=False)
          (bn): BatchNorm2d(24, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
    )
  )
)

```

```

        (1): BasicConv2d(
          (conv): Conv2d(24, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
          (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
    (branch4): Sequential(
      (0): MaxPool2d(kernel_size=3, stride=1, padding=1, dilation=1,
ceil_mode=True)
      (1): BasicConv2d(
        (conv): Conv2d(512, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
  (inception4d): Inception(
    (branch1): BasicConv2d(
      (conv): Conv2d(512, 112, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn): BatchNorm2d(112, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch2): Sequential(
      (0): BasicConv2d(
        (conv): Conv2d(512, 144, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(144, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
      )
      (1): BasicConv2d(
        (conv): Conv2d(144, 288, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn): BatchNorm2d(288, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    (branch3): Sequential(
      (0): BasicConv2d(
        (conv): Conv2d(512, 32, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(32, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
      )
      (1): BasicConv2d(
        (conv): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
  )
)

```

```

    )
    )
    (branch4): Sequential(
      (0): MaxPool2d(kernel_size=3, stride=1, padding=1, dilation=1,
ceil_mode=True)
      (1): BasicConv2d(
        (conv): Conv2d(512, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
  )
  (inception4e): Inception(
    (branch1): BasicConv2d(
      (conv): Conv2d(528, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn): BatchNorm2d(256, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch2): Sequential(
      (0): BasicConv2d(
        (conv): Conv2d(528, 160, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(160, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
      )
      (1): BasicConv2d(
        (conv): Conv2d(160, 320, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn): BatchNorm2d(320, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    (branch3): Sequential(
      (0): BasicConv2d(
        (conv): Conv2d(528, 32, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(32, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
      )
      (1): BasicConv2d(
        (conv): Conv2d(32, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn): BatchNorm2d(128, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    (branch4): Sequential(
      (0): MaxPool2d(kernel_size=3, stride=1, padding=1, dilation=1,
ceil_mode=True)

```

```

        (1): BasicConv2d(
          (conv): Conv2d(528, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
          (bn): BatchNorm2d(128, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
    )
    (maxpool4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=True)
    (inception5a): Inception(
      (branch1): BasicConv2d(
        (conv): Conv2d(832, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(256, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
      )
      (branch2): Sequential(
        (0): BasicConv2d(
          (conv): Conv2d(832, 160, kernel_size=(1, 1), stride=(1, 1), bias=False)
          (bn): BatchNorm2d(160, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
        (1): BasicConv2d(
          (conv): Conv2d(160, 320, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
          (bn): BatchNorm2d(320, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
      (branch3): Sequential(
        (0): BasicConv2d(
          (conv): Conv2d(832, 32, kernel_size=(1, 1), stride=(1, 1), bias=False)
          (bn): BatchNorm2d(32, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
        (1): BasicConv2d(
          (conv): Conv2d(32, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
          (bn): BatchNorm2d(128, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
      (branch4): Sequential(
        (0): MaxPool2d(kernel_size=3, stride=1, padding=1, dilation=1,
ceil_mode=True)
        (1): BasicConv2d(
          (conv): Conv2d(832, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
          (bn): BatchNorm2d(128, eps=0.001, momentum=0.1, affine=True,

```

```

track_running_stats=True)
    )
    )
    (inception5b): Inception(
        (branch1): BasicConv2d(
            (conv): Conv2d(832, 384, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (bn): BatchNorm2d(384, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
        (branch2): Sequential(
            (0): BasicConv2d(
                (conv): Conv2d(832, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
                (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
            )
            (1): BasicConv2d(
                (conv): Conv2d(192, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
                (bn): BatchNorm2d(384, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
            )
        )
        (branch3): Sequential(
            (0): BasicConv2d(
                (conv): Conv2d(832, 48, kernel_size=(1, 1), stride=(1, 1), bias=False)
                (bn): BatchNorm2d(48, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
            )
            (1): BasicConv2d(
                (conv): Conv2d(48, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
                (bn): BatchNorm2d(128, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
            )
        )
        (branch4): Sequential(
            (0): MaxPool2d(kernel_size=3, stride=1, padding=1, dilation=1,
ceil_mode=True)
            (1): BasicConv2d(
                (conv): Conv2d(832, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
                (bn): BatchNorm2d(128, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
            )
        )
    )
    (aux1): None

```



```

(aux2): None
(avgpool): AdaptiveAvgPool2d(output_size=(1, 1))
(dropout): Dropout(p=0.2, inplace=False)
(fc): Linear(in_features=1024, out_features=1000, bias=True)
)

```

0.5 Loss and optimizer

```

[19]: criterion = nn.CrossEntropyLoss()
      optimizer = optim.Adam(model.parameters(), lr = 1e-3)

```

```

[20]: device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')

```

```

[21]: num_epochs = 2

```

0.6 training

```

[25]: for epoch in range(num_epochs):
      losses = []

      for batch_idx, (data, targets) in enumerate(train_loader):
          data = data.to(device = device)
          targets = targets.to(device = device)

          #forward
          scores = model(data)
          loss = criterion(scores, targets)

          losses.append(loss.item())

          #backward
          optimizer.zero_grad()
          loss.backward()

          #adam step
          optimizer.step()
      print(f"Cost at epoch {epoch} is {sum(losses)/len(losses)}")

```

Cost at epoch 0 is 0.026988541723408633

Cost at epoch 1 is 0.007583270934344961

0.7 check accuracy

```

[30]: def check_accuracy(loader, model):
      num_correct = 0
      num_samples = 0
      model.eval()

```

```

with torch.no_grad():
    for x,y in loader:
        x = x.to(device = device)
        y = y.to(device = device)

        scores = model(x)
        _, predictions = scores.max(1)
        num_correct += (predictions == y).sum()
        num_samples += predictions.size(0)

    print(f"Got {num_correct} / {num_samples} with accuracy_{float(num_correct)/float(num_samples)*100}")
    model.train()

```

```

[31]: print("Checking accuracy on Training Set")
      check_accuracy(train_loader, model)

```

Checking accuracy on Training Set
Got 26369 / 26378 with accuracy 99.96588065812419

```

[32]: print("Checking accuracy on Testing Set")
      check_accuracy(test_loader, model)

```

Checking accuracy on Testing Set
Got 11302 / 11305 with accuracy 99.97346306943831

```

[ ]:

```